



Agent Self-Evolution

Phase 2 Validation Report
Closed-loop-aware deploy gate + tool-side parity

Date: May 25, 2026
Repository: github.com/jramos/agent-self-evolution

Executive Summary

Agent Self-Evolution is a standalone optimization pipeline that uses DSPy and GEPA (Genetic-Pareto Prompt Evolution) to automatically improve an agent's skills, tool descriptions, system prompts, and code through evolutionary search — all via API calls with no GPU training required. Phase 1 shipped the framework's first deploy gate (synthetic-only, paired-bootstrap CI, knee-point selection). Phase 2 makes the gate *behavior-aware* — it can now ship candidates whose synthetic signal is flat or slightly negative when the closed-loop behavioral signal demonstrably improves — and ships *tool-side parity*, so the same gate, audit trail, and automation run against tool descriptions as well as skill files.

This report documents the Phase 2 validation of the closed-loop-aware deploy gate. We evolved the **weakened-systematic-debugging** skill end-to-end with the gpt-5.4-mini optimizer stack and validated the candidate against a five-task behavioral suite executed end-to-end by the openai/gpt-5-mini validator agent. The synthetic holdout delta came in tiny-and-negative ($0.976 \rightarrow 0.972$, $\Delta -0.004$) — under Phase 1's strict-synthetic gate this candidate would have been rejected. The closed-loop signal told a different story: behavioral pass-rate went from **2/5** to **5/5** (gained **+3** tasks, required ≥ 1). The new gate consulted that signal directly and decided **DEPLOYED** via the closed-loop behavioral signal; the synthetic sanity check passed (delta within the ± 0.05 tolerance, so the synthetic regression is confirmed within noise).

KEY RESULT — weakened-systematic-debugging (skill-side deploy via CL-aware gate)

Synthetic holdout (n=20): 0.976 → 0.972 ($\Delta -0.004$)

Closed-loop tasks: 2/5 → 5/5 (+3, required ≥ 1)

Synthetic sanity check: passed (Δ within ± 0.05)

Decision: **DEPLOYED** via the closed-loop behavioral signal

Background

Agent Self-Evolution targets the instructions layer of an LLM agent — skill files, tool descriptions, and system prompts — and evolves the text via API-only evolutionary search. The framework was originally built for Hermes Agent (Nous Research) but a pluggable SkillSource protocol now discovers artifacts in the Hermes Agent layout, the Claude Code plugin cache, or any flat local directory. An agent's behavior is governed by three layers:

Layer	What It Is	How It's Currently Improved
-------	------------	-----------------------------

Model Weights	The underlying LLM (Claude, GPT, etc.)	RL training (Tinker-Atropos)
Instructions	Skills, system prompts, tool descriptions	Manual authoring (static)
Tool Code	Python implementations of each tool	Manual development

Phase 1 validated the framework end-to-end on the instructions layer with a synthetic-only deploy gate. Phase 2 extends that gate in two ways. First, the gate is now **behavior-aware**: alongside the synthetic LLM-as-judge holdout, every candidate is exercised on a closed-loop behavioral suite — real tasks scored by a validator agent — and the gate can deploy on the behavioral signal directly when the synthetic signal is flat. Synthetic eval can saturate (every candidate scores near 1.0) or drift; the closed-loop signal answers the question that actually matters at deploy time: *does this candidate help the agent succeed at real tasks?* Second, the same gate, runner, and automation now apply to **tool descriptions** via `evolve_tool.py`, achieving full parity with the skill-side `evolve_skill.py` pipeline.

Approach: Evolutionary Skill Optimization

Three Optimization Engines

Engine	What It Optimizes	License	Role
DSPy + GEPA	Skills, prompts, tool descriptions	MIT	Primary (validated)
DSPy MIPROv2	Few-shot examples, instruction text	MIT	Fallback optimizer
Darwinian Evolver	Code files, algorithms	AGPL v3	Code evolution (Phase 4)

GEPA (Genetic-Pareto Prompt Evolution) is the star engine — an ICLR 2026 Oral paper from Stanford/UC Berkeley. Unlike traditional evolutionary search that only sees pass/fail scores, GEPA reads full execution traces to understand *why* things failed, then proposes targeted mutations. It outperforms reinforcement learning (GRPO) by +6% with 35x fewer rollouts, and outperforms DSPy's previous best optimizer (MIPROv2) by +10%. It works with as few as 3 training examples. Phase 2 ships two refinements to how GEPA is driven: a **saturation pre-flight** that refuses to spend budget on baselines with no headroom, and an **improvement-or-equal acceptance criterion** that halves the false-rejection rate at the noise floor.

The Optimization Pipeline

1. **Saturation pre-flight** — Score the baseline on a sample of synthetic + closed-loop examples; refuse to run if the baseline lands in the `no_headroom`, `uniform_failure`, or `weak_signal` band (cost-saving guard added this cycle)

- 2. **Discover and load artifact** — Resolve the skill (or tool) via the SkillSource / ToolSource protocol, parse YAML frontmatter and body
- 3. **Generate eval dataset** — An LLM reads the artifact and synthesizes (task, expected_behavior) pairs, then splits into train / val / holdout
- 4. **Wrap as DSPy module** — The artifact text becomes a parameterized DSPy module where the instructions are the optimizable parameter
- 5. **Run optimizer** — DSPy GEPA evolves the instructions with improvement-or-equal acceptance, scored by an LLM-as-judge with a structured rubric
- 6. **Knee-point selection** — Among candidates within ϵ of the val-best, pick the highest-val candidate (smallest body as tiebreak)
- 7. **Dual-signal deploy gate** — Score the candidate on the synthetic holdout (paired-bootstrap CI) AND execute it on a closed-loop behavioral suite; the gate consults a decision_signal field and may deploy via the closed-loop path when synthetic is flat but behavior gains $\geq$ required threshold
- 8. **Report** — Structured gate_decision.json (v5 schema, with CL-primary fields), before/after artifacts, full LM trace log, opt-in --create-pr automation

The **saturation pre-flight** is the new cost-control story for Phase 2. Synthetic LLM-as-judge fitness saturates aggressively at this validator tier: in preparing this report we ran the pre-flight against three candidate headline artifacts and two of them — search_files (a tool description) and a deliberately-weakened write_file suite — were correctly refused as saturated against the validator's gpt-5.4-mini tier. The pre-flight kept us from spending GEPA budget on baselines GEPA cannot beat. The headline run reported here (**weakened-systematic-debugging**) cleared the pre-flight in the weak_signal band — the band the gate was redesigned for — and consumed **\$1.27** across 689 LM calls in ~88 minutes. **Tool-side parity** is the second Phase 2 deliverable: evolve_tool.py ships the same GEPA runner, dataset builder, quality-gate, audit trail, and opt-in PR automation as evolve_skill.py. The headline result lands skill-side because the closed-loop suites we have for tool-side surfaces are saturated for our current validator tier — the deliverable is full parity, ready for harder tool-side eval surfaces.

Phase 2 Experiment

Configuration

Parameter	Value
Target Skill	weakened-systematic-debugging (weakened systematic-debugging — deliberately-weakened baseline that lands in the weak_signal saturation band, exercising the CL-aware deploy path)
Baseline Size	1,987 characters

Optimizer LM	openai/gpt-5.4-mini
Reflection LM (GEPA)	openai/gpt-5.4-mini
Eval / Judge LM	openai/gpt-5.4-mini
Optimizer	DSPy GEPA (light budget; improvement-or-equal acceptance)
Synthetic Eval Set	53 examples (18 train / 15 val / 20 holdout)
Total Optimization Time	5,305 seconds (~88 minutes)
Total LM Calls	689 (from metrics.json per-model summary)
Total Cost (USD)	\$1.27
Quality Gate	dual-signal — synthetic holdout (paired-bootstrap CI) + closed-loop behavioral suite (CL-primary on tie)
Knee-point Strategy	val-best (default; smallest available via --knee-point-strategy)
Closed-loop Validator	openai/gpt-5-mini
Closed-loop Suite	5 tasks (behavioral benchmark, scored end-to-end)

Evaluation Dataset

The evaluation dataset was synthetically generated by openai/gpt-5.4-mini. Given the full weakened-systematic-debugging SKILL.md text, the model produced 53 realistic test cases with rubric-based expected behaviors, then split them into train / val / holdout per the framework's configured ratios. The closed-loop validation suite is a separate five-task `systematic_debugging.jsonl` behavioral benchmark executed end-to-end by the validator agent — those tasks are not part of the synthetic split and are evaluated only on the baseline and the final knee-point candidate. Examples drawn from the synthetic train split:

Task Input	Expected Behavior (Rubric)
The function <code>`average_nonzero(nums)`</code> in <code>`analysis.py`</code> divides by <code>`len(nums)`</code> even though it should ignore zeros. Diagnose the bug and specify the fix.	State that the denominator is wrong because zeros are included, explain current vs intended behavior, and propose filtering zeros before computing the mean.
The function <code>`median(values)`</code> in <code>`stats.py`</code> sorts the input list in place. A test expects the original list to remain unchanged. Diagnose the bug.	Explain that in-place sort mutates caller data, identify the sort line, state intended behavior should preserve input sequence, and propose sorting a copy instead.
In <code>`prime.py`</code> , <code>`is_prime(n)`</code> returns True for 1. Diagnose the bug and specify the exact fix.	Identify missing guard for $n < 2$, explain current behavior incorrectly treats 1 as prime, intended behavior excludes 0 and 1, and propose adding an early return.

Fitness Function

Synthetic fitness is measured by an LLM-as-judge (gpt-5.4-mini) that scores each candidate output along three rubric dimensions. The composite score is a weighted combination with a

length-penalty term that discourages runaway expansion:

$$\text{composite} = 0.5 \cdot \text{correctness} + 0.3 \cdot \text{procedure_following} + 0.2 \cdot \text{conciseness} - \text{length_penalty}$$

The judge also returns a free-text feedback string that GEPA's reflection LM consumes to propose targeted instruction-text mutations on the next iteration — this trace-aware loop is the core of GEPA's sample efficiency. Phase 2 adds a second, independent fitness signal at gate time: **closed-loop behavioral pass-rate** on a small held-out task suite executed by a validator agent. The gate consults both signals — synthetic for sanity (± 0.05 tolerance), closed-loop for the deploy decision when synthetic is flat.

Results

Metric	Baseline	Evolved (knee-point pick)	Δ
Body size (chars)	1,987	2,967	+49.3%
Avg holdout score (n=20)	0.976	0.972	-0.004
Bootstrap mean diff	—	-0.004	—
Bootstrap 90% CI lower	—	-0.028	—
Bootstrap 90% CI upper	—	+0.012	—
Closed-loop tasks (n=5)	2/5	5/5	+3 (req ≥ 1)
Decision	—	DEPLOYED	via closed-loop

The evolved weakened-systematic-debugging skill grew **+49.3%** (1,987 $\rightarrow$ 2,967 chars) and the synthetic holdout score moved **-0.004** (0.976 $\rightarrow$ 0.972) on n=20 examples. The synthetic delta is tiny-and-negative — under Phase 1's no-regression rule, this candidate would have been rejected. Phase 2's behavior-aware gate looked at the closed-loop signal instead: behavioral pass-rate went from **2/5** tasks to **5/5** (gained **+3**, required ≥ 1). The synthetic sanity check passed (delta inside the ± 0.05 noise envelope), so the synthetic regression is statistically indistinguishable from noise while the behavioral improvement is large and concrete. **Decision: DEPLOYED** via the closed-loop behavioral signal. This is the textbook case the Phase 2 gate was redesigned for.

How the Result Was Produced

GEPA evolves skill instructions through a reflective loop; Phase 2's gate then reads two independent signals at decision time:

1. Run candidate skill instruction text on training examples; the judge scores each output and emits free-text feedback
2. Reflection LM reads the execution traces + feedback and proposes a targeted mutation of the instruction text (Phase 2: improvement-or-equal acceptance keeps near-ties in the

population)

3. Score the mutated candidate on the validation set (15 examples); track every candidate's per-example Pareto front
4. After GEPA's light-budget search, freeze the candidate population; knee-point selection picks candidate 1 (val=0.987, rank 1 of 2 in the ϵ -band, 2,753 body chars) and is also the candidate GEPA's own default selector would have chosen — the val-best and GEPA-default converged here
5. **Dual-signal gate** — Score the knee-point pick on the 20-example synthetic holdout (paired-bootstrap CI on per-example diffs) AND execute it on the closed-loop behavioral suite (5 tasks, scored by the openai/gpt-5-mini validator). The gate sets `decision_signal` based on which signal carries the decision; on this run synthetic was flat-within-tolerance and closed-loop gained 3 tasks, so the gate deployed via the closed-loop path.

Three Phase 2 design choices made this outcome possible: (a) the dual-signal gate deploys on either signal, so a flat synthetic doesn't veto a real behavioral gain; (b) the synthetic sanity check still guards against unambiguous synthetic regressions (± 0.05 tolerance); (c) the saturation pre-flight refused two of three candidate headline artifacts before this run, which is honest evidence that the current validator tier saturates aggressively on our existing eval surfaces — the Phase 2 gate exists specifically to recover deploy decisions on the artifacts that do clear the pre-flight in the *weak_signal* band. The CL-aware deploy gate arc is supported by a broader May-cycle calibration campaign across nano-pdf, apple-notes, polymarket, and huggingface-hub (reports/calibration_findings.md) — that campaign contributed the improvement-or-equal acceptance default, retired the knee-point ϵ selector as a no-op on val-best, and recommended the non-inferiority tolerance sweet spot used by the synthetic sanity check.

Safety and Guardrails

Every evolved variant must pass all of the following constraints before deployment:

Constraint	Enforcement	Status
Self-evolution test suite	1,166 pytest tests pass on the optimizer itself	Implemented
Static size limits	Skills ≤ 15 KB, tool descs ≤ 500 chars (configurable)	Implemented
Absolute char ceiling	Hard cap on evolved artifact size (default 5,000)	Implemented
Growth-quality gate	Required improvement scales linearly with growth %	Implemented
Paired-bootstrap CI	90% CI on per-example holdout diffs gates deploy	Implemented

Knee-point selection	Smallest candidate within ϵ of val-best	Implemented
Structural integrity	Valid YAML frontmatter required	Implemented
Deployment via PR	Human review required, never auto-merge	By design
CL-aware deploy gate	Deploys via closed-loop signal when CL gain $\geq$ required AND synthetic delta within ± 0.05	Implemented
Saturation pre-flight	Refuses to spend budget on no_headroom / uniform_failure / weak_signal baselines	Implemented
GEPA improvement-or-equal	Candidates accepted on $\geq$ (not strict $>$); halves false-rejection at the noise floor	Implemented
PR automation audit trail	Opt-in --create-pr opens draft PR; pr_created field logs branch + SHA + URL	Implemented
Benchmark regression	TBLite / skill-specific harness must hold	Planned

Source skill and tool repositories are never modified directly. All evolution output (evolved artifacts, gate decisions, run logs, closed-loop validation transcripts) is written under the framework's local output/ directory, and improvements are proposed as draft pull requests against the source repo for human review.

Roadmap

Phase	Target	Engine	Timeline	Status
Phase 1	Skill files (SKILL.md)	DSPy + GEPA	3-4 weeks	Validated ✓
Phase 2	Tool descriptions	DSPy + GEPA	2-3 weeks	Validated ✓
Phase 3	System prompt sections	DSPy + GEPA	2-3 weeks	Planned
Phase 4	Tool implementation code	Darwinian Evolver	3-4 weeks	Planned
Phase 5	Continuous improvement	Automated pipeline	2 weeks	Planned

Each phase must demonstrate measurable improvement and pass benchmark regression gates before proceeding. Phase 2 ships the deploy gate's behavioral awareness AND tool-side parity — the gate is no longer hostage to synthetic-eval saturation, and the same automation that gated skill evolution now gates tool-description evolution. Phase 3 (system-prompt sections) is the next surface; Phase 5 (continuous improvement) closes the loop with automated cron-driven optimization.

Immediate Next Steps

1. **Phase 3 scoping** — System-prompt sections as the next evolvable surface. The instructions layer remains the target; system prompts complete the trio of skills / tools / system prompts and are the highest-leverage instructions surface for most agents.
2. **Eval-surface hardening** — Develop harder closed-loop suites and synthetic generators so the saturation pre-flight is less often the dominant outcome at the gpt-5.4-mini validator tier. The Phase 2 gate works; the bottleneck is now the eval surfaces feeding it.
3. **Cross-tool portfolio campaign** — Run the CL-aware deploy gate against the realistic manifest's tool descriptions to surface which tool-side artifacts have headroom under the current validator tier — analogous to the May skill-side calibration campaign that produced the improvement-or-equal acceptance default.
4. **Phase 5 — continuous improvement** — Cron-driven optimization with budget gates, alerting, and an opt-in PR-automation queue. The `--create-pr` primitive shipped in Phase 2 is the prerequisite; Phase 5 wires it into a scheduled loop with backstop budgets.
5. **Documentation polish** — Link the Phase 2 report from the README roadmap section and the framework's headline narrative; refresh the architecture diagram to reflect the dual-signal gate.